

Weak Memory Concurrency-II

Soham Chakraborty

27.02.2023

Semantics

Reasoning about Properties

- Coherence
- Atomicity

Concurrency and compilation

Semantics Models for Concurrency

Transformation+Interleaving

Operational semantics

Axiomatic semantics

Program is a set of *consistent* execution

Execution = events + relations

- events represent accesses on shared memory locations or fences
- Events are related by binary Relations

Axioms enforce constraints on an execution

Memory model = set of axioms

Given a memory consistency model M and an execution X , X is M -consistent if X satisfies all the axioms of M

event =

$\langle \text{unique identifier, thread id, operation, order, location, value} \rangle$

Examples

A release write access writes value v on location X

- $\text{St}_{\text{rel}}(X, v)$

A acquire read access reads value v from location X

- $\text{Ld}_{\text{acq}}(X, v)$

Each relation relate two events

Example

Relation Program-Order (po): captures syntactic order of the events

Relation reads-from (rf):

If a read r reads-from a write w then $(w, r) \in rf$

Constraints:

- w and r access same memory locations
- Written value of w = read value of r

Modification order (mo): orders write events that access the same memory location

Execution Graph

Execution is represented as a graph

Nodes are events and edges are relations

Axioms are (usually) acyclicity conditions

Example: No read access reads from a later write

More Relations

Relations can be primitive or derived from other relations

From-read(fr): orders a read to a later write on the same location

Derived from rf and **mo**

$$\text{fr} \triangleq \text{rf}^{-1}; \text{mo}$$

Sequential Consistency

All shared memory accesses are ordered

$(po \cup rf \cup \text{mo} \cup fr)$ are acyclic.

Example: Forbid weak execution(s)

$$\begin{array}{l} X = Y = 0; \\ X = 1; \parallel Y = 1; \\ a = Y; \parallel b = X; \end{array}$$

Outcome $a = b = 0$ are disallowed

Coherence Property

SC per location

A new relation:

poloc: po between events that access the same memory location

$\text{poloc} = \{(a, b) \mid \text{po}(a, b) \wedge a.\text{location} = b.\text{location}\}$

Coherence: $(\text{poloc} \cup \text{rf} \cup \text{fr} \cup \text{mo})$ is acyclic

Examples: coherence violations

$X = 1;$

$X = 2;$

$a = X; \text{ // } 1$

$a = X; \text{ // } 1$

$X = 1;$

$X = 1; \parallel a = X; \text{ // } 2$
 $X = 2; \parallel b = X; \text{ // } 1$

$X = 1; \parallel X = 2;$
 $a = X; \text{ // } 2 \parallel b = X; \text{ // } 1$

Atomicity: $\text{rmw} \cap (\text{fr}; \text{mo}) = \emptyset$

Examples: Atomicity violations

$$\begin{array}{c} X = 0; \\ \text{CAS}(X, 0, 1); \parallel \text{CAS}(X, 0, 1); \end{array}$$

Both CAS operations cannot be successful

Release-Acquire Consistency

All writes are release and all reads are acquire accesses

Follows (sc-per-loc) and (atomicity)

Reordering restrictions: WW, RR, RW (same as TSO)

Example: Allowed behaviors (same as TSO in these programs)

```
X = Y = 0;  
  
X = 1; || a = Y; // 1  
Y = 1; || b = X; // 0
```

```
X = Y = 0;  
  
X = 1; || Y = 1;  
a = Y; // 0 || b = X; // 0
```

Release-Acquire Consistency

Reordering restrictions: WW, RR, RW (same as TSO)

Allows non-multicopy atomicity unlike TSO

Example:

$$\begin{array}{c} X = Y = 0; \\ X = 1; \parallel \begin{array}{l} a = X; \\ b = Y; \end{array} \parallel \begin{array}{l} c = Y; \\ d = X; \end{array} \parallel Y = 1; \end{array}$$

Outcome $a = c = 1, b = d = 0$ is allowed in RA but not in TSO

Release-Acquire Consistency

New relation

Happens-before: $hb \triangleq (po \cup rf)^+$

Identify the hb relations in earlier examples

Axioms

$(hb \cup rf \cup fr \cup \text{mo})$ is acyclic

(sc-per-loc)

$\text{rmw} \cap (fr; \text{mo}) = \emptyset$

(atomicity)

Non-atomic accesses

Relaxed accesses

Acquire, release accesses and fences

SC accesses and SC fence

Formal model is known as C11

Synchronization relation is established by

- Release acquire accesses
- Relaxed accesses and fences

Happens-before: $hb = (po \cup sw)^+$

Data Race

a and b is a data race (on non-atomics) when

- a and b are concurrent (not related by hb)
- Access same memory location
- Atleast one of a or b is non-atomics

A consistent execution with data race on non-atomic

\implies the behavior of the program is undefined

$$X = 0$$

$$X_{\text{na}} = 1; \parallel a = X_{\text{acq}};$$

$a > 1$ is possible in C/C++

OOTA: out-of-thin-air behavior

$$\begin{array}{c}
 X = Y = 0 \\
 a = X; \quad \parallel \quad b = Y; \quad (\text{LBDep}) \\
 \text{if}(a == 1) \parallel \text{if}(b == 1) \\
 Y = 1; \quad \parallel \quad X = 1;
 \end{array}
 \quad \Bigg| \quad
 \begin{array}{c}
 X = Y = 0 \\
 a = X; \parallel b = Y; \quad (\text{LB}) \\
 Y = 1; \parallel X = 1;
 \end{array}$$

Undesirable: C11 allows $a = b = 1$ in both programs

Desirable: Allow $a = b = 1$ in (LB), but forbid in (LBDep)

Correctness of Compilation

Correct transformations in sequential programs may NOT be correct in concurrent programs

Example: Under RA model

$X = Y = 0;$		$X = Y = 0;$
$X = 1; \parallel a = Y;$	\rightsquigarrow	$X = 1; \parallel a = Y;$
$Y = 1; \parallel b = X;$		$Y = 1; \parallel b = X;$
$a = 1, b = 0$ is forbidden		$a = 1, b = 0$ is allowed

Incorrect transformation

For each consistent execution of the target program there exists a corresponding consistent execution of the source program with same outcome.

Proof Strategy

- 1 For each consistent execution X of the target program, define an execution X' of the source program.
- 2 Show that X' is consistent w.r.t the source consistency model
- 3 Show the behavior of X and X' are same.

Transformations:

- Program optimization
- Mapping to architectures

Correct Program Transformations

Reordering transformations: $a; b; \rightsquigarrow b; a;$

Elimination transformations:

- $t = X; t' = X; \rightsquigarrow t = X; t' = t;$ (RAR)
- $X = v; t = X; \rightsquigarrow X = v; t = v;$ (RAW)
- $X = v; X = v'; \rightsquigarrow X = v';$ (OW)

Mapping Schemes

Programming languages primitives to architectures

Additional leading and/or trailing fences

Example

<https://www.cl.cam.ac.uk/~pes20/cpp/cpp0xmappings.html>

Main Challenge: Proving mapping correctness

Example Mappings: RA to TSO

$$W_{\text{rel}} \rightsquigarrow W, R_{\text{acq}} \rightsquigarrow R, \text{CAS} \rightsquigarrow \text{CAS}, F \rightsquigarrow F$$

Mapping Correctness: Suppose a program P in RA is mapped to P' in TSO following the above mapping scheme. For each TSO-consistent execution of P' there exists an RA-consistent execution of P having same behavior.

- Behavior: Final values in the shared memory locations

Why different memory models?

Role of compilers

Considerations for a new memory model

Analysis tools

Mathematizing C++ Concurrency.

Mark Batty, Scott Owens, Susmit Sarkar, Peter Sewell, and Tjark Weber.

In POPL 2011

Chapter 2 (Background)

Correct Compilation of Relaxed Memory Concurrency.

Soham Chakraborty.

`http:`

`//plv.mpi-sws.org/soham/thesis/Thesis-Chakraborty.pdf`

`https://www.cse.iitd.ac.in/~soham/COL869/page.html`

Lectures: 30.03.2020 - 24.04.2020