

Concurrency Bugs

Soham Chakraborty

16.02.2022

Order violation

Atomicity violation

Deadlock

Data race

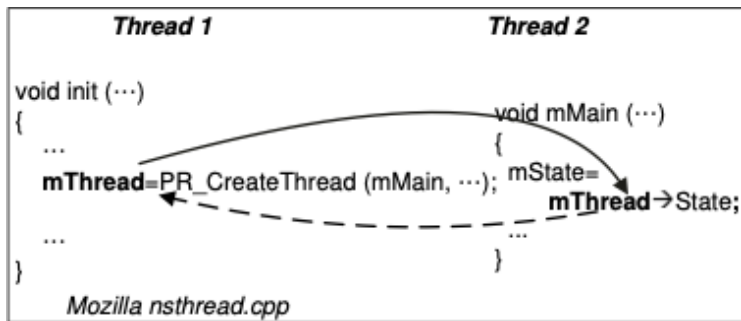
Order Violation

Cause: Programmer assumes certain ordering of events

Order Violation

Cause: Programmer assumes certain ordering of events

Example:




Thread 2 should not deref. **mThread** before Thread 1 initializes it

Order Violation

Cause: Programmer assumes certain ordering of events

Example:




Thread 2 should not deref. **mThread** before Thread 1 initializes it
Pattern:

```
X = 0;  
X = 1; || t = X; // 1
```

Order Violation

Cause: Programmer assumes certain ordering of (W,R) events

Example:




`js_UnpinPinnedAtom` should happen after `js_MarkAtom`.

Order Violation

Cause: Programmer assumes certain ordering of (W,R) events

Example:



`js_UnpinPinnedAtom` should happen after `js_MarkAtom`.

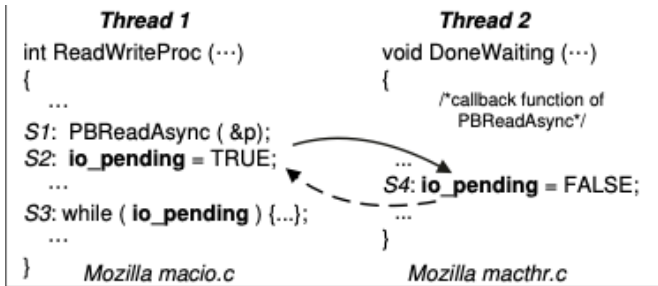
Pattern:

```
X = 0;  
X = 1; || t = X; // 0
```

Order Violation

Cause: Programmer assumes certain ordering of (W,W) events

Example:




Assumption: S1 and S2 execute atomically

Unsafe ordering blocks thread 1

Order Violation

Cause: Programmer assumes certain ordering of (W,W) events

Example:



Assumption: S1 and S2 execute atomically

Unsafe ordering blocks thread 1

Pattern:

```


X = 1;
while(X == 1) ; // 0 || X = 0;

```

Atomicity Violation

Cause: Programmer assumes atomicity of certain code regions

Example:




Assumption: S1;S2 are executed atomically

S2 access NULL value

Atomicity Violation

Cause: Programmer assumes atomicity of certain code regions

Example:



Assumption: S1;S2 are executed atomically

S2 access NULL value

Pattern:


$$\begin{array}{l} X = 0; \\ a = X; \\ b = X; \end{array} \parallel \begin{array}{l} X = 1; \end{array}$$

Desired: $a = b = 0$ or $a = b = 1$

Multi-Variable Atomicity Bugs

Cause: variables are semantically connected which is violated

Example:




Assumption: `mOffset` and `mLength` are updated atomically wrt thread 1

Lack of synchronization \Rightarrow thread 1 read inconsistent value

Multi-Variable Atomicity Bugs

Cause: variables are semantically connected which is violated

Example:



Assumption: `mOffset` and `mLength` are updated atomically wrt thread 1

Lack of synchronization \Rightarrow thread 1 read inconsistent value

$$Y = Z = 0;$$
$$t = X[Y + Z]; \parallel Y = 1; Z = 1;$$

Desired: access `X[0]` or `X[2]`

Timing Bugs

Cause: Programmer assumes the tasks would complete within certain time period

Example:

<i>Thread 1</i>	<i>Thread 2 ... Thread n</i>	<i>Monitor thread</i>
<pre>void buf_flush_try_page() { ... rw_lock(&lock); }</pre>	<pre>rw_lock(&lock);</pre>	<pre>void error_monitor_thread() { if(lock_wait_time[i] > fatal_timeout) assert(0, "We crash the server; It seems to be hung."); }</pre>
<i>MySQL buf0flu.c</i>		<i>MySQL srv0srv.c</i>

Assumption: n tasks would complete before *fatal_timeout*

Crash the server

Fix Strategies

Understand the semantics

Add/modify locks

Add/modify synchronizations

Revisit the examples

Deadlock

A thread holds a lock and wait for another lock held by another thread and vice versa

<i>lock(m₁);</i>	<i>lock(m₂);</i>
<i>lock(m₂);</i>	<i>lock(m₁);</i>
<i>...</i>	<i>...</i>
<i>unlock(m₂);</i>	<i>unlock(m₁);</i>
<i>unlock(m₁);</i>	<i>unlock(m₂);</i>

Deadlock: Another Scenario

Another challenge: encapsulation

```
Vector v1, v2;  
v1.AddAll(v2); || v2.AddAll(v1);
```

Conditions for Deadlock

All conditions must hold:

- **Mutual exclusion:** Threads claim exclusive control of resources (e.g. lock) that they require.
- **Hold-and-wait:** Threads hold allocated resources while waiting for additional resources
- **No preemption:** Held resources cannot be forcibly removed from threads
- **Circular wait:** There exists a circular chain of threads where each thread holds a resource that are being requested by the next thread in the chain.

Prevent circular wait Programming convention: total ordering on acquiring lock

- Prone to mistakes

Prevent hold-and-wait Acquire all locks at once

- Decreases concurrency significantly

Prevent no-preemption

Hold locks only when all the locks are available

Challenge: encapsulation prevents the 'top' loop implementation

```
top :  
  lock(L1);  
  if(trylock(L2) == -1) {  
    unlock(L1);  
    goto top;  
  }
```

Prevent no-preemption

Hold locks only when all the locks are available

Challenge: encapsulation prevents the 'top' loop implementation

<pre><i>top</i> : lock(L1); if(trylock(L2) == -1) { unlock(L1); goto top; }</pre>		<pre><i>top</i> : lock(L2); if(trylock(L1) == -1) { unlock(L2); goto top; }</pre>
---	--	---

Prevent no-preemption

Hold locks only when all the locks are available

Challenge: encapsulation prevents the 'top' loop implementation

<pre><i>top</i> : lock(L1); if(trylock(L2) == -1) { unlock(L1); goto top; }</pre>		<pre><i>top</i> : lock(L2); if(trylock(L1) == -1) { unlock(L2); goto top; }</pre>
---	--	---

Problem: Livelock

Prevent circular wait Total ordering on acquiring lock

- Prone to mistakes

Prevent hold-and-wait Acquire all locks at once

- Decreases concurrency significantly

Prevent no-preemption

- Problem: Livelock
- Challenge: encapsulation prevents the 'top' loop implementation


No mutual-exclusion

- Lock free programming

Deadlock Avoidance

Schedule threads that access same resources

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no



Deadlock Recovery


Deadlock detector automatically detect deadlock

If deadlock is detected; restart system

Event a and b is in data race if:

- a and b are concurrent/in conflict
- a and b access same location
- At least one of a and b is a write

Concurrent Accesses




Concurrent: (e_1, e_2) , (e_2, e_3)

e_3 happens-before e_1

- $end(e_3) \rightarrow start(e_1)$

Happens-Before

concurrent/conflict \Rightarrow Not in happens-before (HB) order



Execution 1: No data race

Execution 2: data race on `x`

Lockset algorithm

Let $locks_held(t)$ be the set of locks held by thread t .

For each v , initialize $C(v)$ to the set of all locks.

On each access to v by thread t ,

 set $C(v) := C(v) \cap locks_held(t)$;

 if $C(v) = \{ \}$, then issue a warning.

Lockset algorithm

Let $locks_held(t)$ be the set of locks held by thread t .

For each v , initialize $C(v)$ to the set of all locks.

On each access to v by thread t ,
set $C(v) := C(v) \cap locks_held(t)$;
if $C(v) = \{ \}$, then issue a warning.

Example:

<i>Program</i>	<i>locks_held</i>	<i>C(v)</i>
	{}	{mu1, mu2}
lock(mu1);	{mu1}	
v := v+1;		{mu1}
unlock(mu1);	{}	
lock(mu2);	{mu2}	
v := v+1;		{}
unlock(mu2);	{}	

Learning from Mistakes – A Comprehensive Study on Real World Concurrency Bug Characteristics.

Shan Lu, Soyeon Park, Eunsoo Seo and Yuanyuan Zhou
ASPLOS 2008.

Common Concurrency Problems (chapter 32)

Operating Systems: Three Easy Pieces

Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau

<https://pages.cs.wisc.edu/remzi/OSTEP/threads-bugs.pdf>

Race Detection Techniques

Christoph von Praun

https://doi.org/10.1007/978-0-387-09766-4_38

Eraser: A Dynamic Data Race Detector for Multithreaded Programs

Stefan Savage, Michael Burrows, Greg Nelson, Patrick Sobalvarro, Thomas Anderson. ACM TOCS 1997.