

Introduction to Distributed Systems

CS4405 – Analysis of Concurrent and Distributed Programs


Burcu Kulahcioglu Ozkan




Shared-memory vs distributed systems



What concurrency problems do you think are common or different in these systems?



Shared-memory





Distributed



What is a distributed system?

- “. . . a system in which the failure of a computer you didn’t even know existed can render your own computer unusable.”
— Leslie Lamport
- The computers in the system:
 - Are connected over network
 - Communicate by exchanging messages



Why distributed systems?

Inherent Distribution:


- Information dissemination (publishers/subscribers)
- Distributed process control
- Cooperative work (different nodes on a network read/write)
- Distributed storage

Distribution as an Artifact:

- Performance
- Scalability (data, geographical, functional)
 - Moore's law: The number of transistors *on a single chip* doubles about every two year.
 - The advancement has slowed since around 2010.
 - Distribution provides massive performance.
- Availability
- Fault tolerance




A simple distributed system example



Fallacies of distributed systems


By [Peter Deutsch](#)

- [The network is reliable](#)
- Latency is zero
- Bandwidth is infinite
- The network is secure
- Topology does not change
- Transport cost is zero
- The network is homogeneous





Distributed system abstractions

- Abstracting the underlying physical system into a system model:
 - Processes and messages
 - Communication network
- We start with a primitive link abstraction (e.g. from point A to point B), and layer on abstractions to build stronger guarantees.
- Abstractions for common interaction patterns (incremental):
 - Process identities (processes agree on who they are)
 - Consensus (processes agree on a common plan)
 - Atomic commitment (processes take a step only if all processes agree)
 - Total order broadcast (processes agree on a total order of actions)



Processes and messages

- Processes: “units that are able to perform computations in a distributed system”
- A distributed system consists of a collection of automata, one per process
- The execution of a distributed algorithm is represented by a sequence of steps executed by the processes:
 - Receiving a message
 - Executing a local computation
 - Sending a message to a process



Abstracting Processes: Process failures

- Crash fault
- Omission fault
- Crash-recovery fault
- Arbitrary/Byzantine fault




System synchrony model: Synchronous

- Process execution speeds or message delivery times are bounded.


In a synchronous system, we can have:

- Timed failure detection
- Time based coordination
- Worst-case performance



System synchrony model: Asynchronous

- No assumptions about process execution or message delivery times are made
- Upon waiting for a response to a requests, it is **not possible to distinguish** whether:
 - the request was lost
 - the remote node is down
 - the response was lost



System synchrony model: Partially-synchronous

The retransmission of the messages may help ensure the reliability of the communication links but introduce unpredictable delays.

Set **timeouts** and retry the request until it succeeds

In this sense, practical systems are partially synchronous:

- **Partially-synchronous system:**
 - There exist upper bounds on the network delay but the programmer does not know them.
 - Eventually synchronous: The network may behave asynchronously for some duration, but eventually behaves synchronously.

