


Time, Order, and Causality in Distributed Systems


CS4405 – Analysis of Concurrent and Distributed Programs

Burcu Kulahcioglu Ozkan



Events in distributed systems

- Processes operate on their local memory and communicate by exchanging messages:
 - A process performs some local computation
 - A process sends a message
 - A process receives a message
- 
- The diagram shows a central server icon (a tall grey rectangle with a blue cylinder on top) connected by blue lines to three other server icons. Each connection is accompanied by a blue envelope icon, representing a message being sent or received. This illustrates how processes (represented by the server icons) communicate via messages.



Time and order of events in distributed systems



Why do we need to order the events?

- Encoding history (“happens before” relationships)
- Transactions in a database
- Consistency of distributed data
- Debugging (finding the root cause of a bug)
- . . .



Reminder: Partial vs Total order

Strict partial order:

- **Irreflexivity:** $\forall a. \neg a < a$ (items not comparable with self)
- **Transitivity:** if $a \leq b$ and $b \leq c$ then $a \leq c$
- **Antisymmetry:** if $a \leq b$ and $b \leq a$ then $a = b$


Strict total order:

- An additional property: $\forall a, b, a \leq b \vee b \leq a \vee a = b$




Time in centralized vs distributed systems

- Centralized systems: System calls to kernel, monotonically increasing time values.
- Distributed systems: Achieving agreement on time is not trivial!



Logical time

- Idea: Instead of using the precise clock time, capture the events relationship between a pair of events
- Based on **causality**: If some event possibly causes another event, then the first event **happened-before** the other




Operating
Systems

R. Stockton Gaines
Editor

Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport
Massachusetts Computer Associates, Inc.

CACM, 1978




Happens-before relation between events

Happens-before relation captures dependencies between events:

- If a and b are events in the same node, and a occurs before b , then $a \rightarrow b$
- If a is the event of sending a message and b is the event of receiving that message, then $a \rightarrow b$
- The relation is transitive.

It is a strict partial order: it is irreflexive, antisymmetric and transitive.

Two events not related to happened-before are **concurrent**.




Lamport timestamps

Lamport introduced the eponymous logical timestamps in 1978:


- Each individual process p maintains a counter: $LT(p)$.
- When a process p performs an action, it increments $LT(p)$.
- When a process p sends a message, it includes $LT(p)$ in the message.
- When a process p receives a message from a process q , that message includes the value of $LT(q)$; p updates its $LT(p)$ to the $\max(LT(p), LT(q)) + 1$

For two events a and b , if $a \rightarrow b$, then $LT(a) < LT(b)$.



Lamport timestamps

For two events a and b , if $a \rightarrow b$, then $LT(a) < LT(b)$.



If $LT(a) < LT(b)$, then it does not mean that $a \rightarrow b$.



Why is the LT invariant not symmetric?

Another example scenario with 4 nodes that exchange events:

Initial state of timestamps: $[A(0), B(0), C(0), D(0)]$

E1. A sends to C : $[A(1), B(0), C(0), D(0)]$

E2. C receives from A : $[A(1), B(0), C(2), D(0)]$

E3. C sends to A : $[A(1), B(0), C(3), D(0)]$

E4. A receives from C : $[A(4), B(0), C(3), D(0)]$

E5. B sends to D : $[A(4), B(1), C(3), D(0)]$

E6. D receives from B : $[A(4), B(1), C(3), D(2)]$

At this point, $LT(E6) < LT(E4)$, but it does not mean that $E6 \rightarrow E4$!

Events 4 and 6 are independent.



Vector Clocks

Vector clocks can maintain causal order.

On a system with N nodes, each node i maintains a vector V_i of size N .

- $V_i[i]$ is the number of events that occurred at node i
- $V_i[j]$ is the number of events that node i knows occurred at node j

All nodes vector clocks start at $[0, \dots, 0]$

They are updated as follows:

- Local events increment $V_i[i]$
- When i sends a message to j , it includes V_i
- When j receives V_i , it updates all elements of V_j to $V_j[a] = \max(V_i[a], V_j[a])$



Vector clocks

Initial state of timestamps: $[A(0, 0, 0, 0), B(0, 0, 0, 0), C(0, 0, 0, 0), D(0, 0, 0, 0)]$

E1. A sends to C : $[A(1, 0, 0, 0), B(0, 0, 0, 0), C(0, 0, 0, 0), D(0, 0, 0, 0)]$

E2. C receives from A : $[A(1, 0, 0, 0), B(0, 0, 0, 0), C(1, 0, 1, 0), D(0, 0, 0, 0)]$

E3. C sends to A : $[A(1, 0, 0, 0), B(0, 0, 0, 0), C(1, 0, 2, 0), D(0, 0, 0, 0)]$

E4. A receives from C : $[A(2, 0, 2, 0), B(0, 0, 0, 0), C(1, 0, 2, 0), D(0, 0, 0, 0)]$

E5. B sends to D : $[A(2, 0, 2, 0), B(0, 1, 0, 0), C(1, 0, 2, 0), D(0, 0, 0, 0)]$

E6. D receives from B : $[A(2, 0, 2, 0), B(0, 1, 0, 0), C(1, 0, 2, 0), D(0, 1, 0, 1)]$



Vector clock guarantees

- Comparing vector clocks: Given V_i and V_j :
 - $V_i = V_j$ iff $V_i[k] = V_j[k]$ for all k
 - $V_i < V_j$ iff $V_i[k] \leq V_j[k]$ for all k and $V_i \neq V_j$
 - (Concurrency) $V_i \parallel V_j$ otherwise
- For two events a and b and their vector clocks $V(a)$ and $V(b)$:
 - if $a \rightarrow b$, then $V(a) < V(b)$
 - if $V(a) < V(b)$, then $a \rightarrow b$

Vector clocks are expensive to maintain: they require $O(n)$ timestamps to be exchanged with each communication.

- However, we cannot do better than $O(n)$



Causally dependent events



Why compute causal dependency between events?

